

A Beam-space Selective State-Space World Model for MIMO-OFDM Channel Prediction and Estimation

Abstract

We present a *world model* for MIMO-OFDM wireless channels: an action-conditioned model of channel dynamics that, given a short history of channels and the receiver-motion action, predicts the future channel and supplies a temporal prior for channel estimation. The model operates in the **beam-space** (angle–delay) domain, where the channel is sparse and a compact latent is reconstruction-faithful; a Mamba-style *selective state-space* (SSM) backbone rolls the latent forward under the action, and a decoder reconstructs the future channel. Trained end-to-end by Distributed Data Parallel across $8\times$ A100 GPUs on 4.8×10^4 ray-traced Sionna RT 8×4 MIMO channel sequences spanning six city scenes, the world model is the best *channel predictor* at every horizon $k=1,\dots,6$, beating persistence and a linear autoregressive baseline with a margin that widens with horizon, both in-distribution and on held-out scenes. For *channel estimation* we report the complete picture, honestly: the world-model prior is decisively useful when the observation is degraded but the history is informative — sparse pilots at moderate SNR — while a lean per-snapshot convolutional estimator is stronger under dense pilots and at SNR extremes. We unify the two into a single framework with a physically-motivated, CSI-adaptive router that matches or beats every baseline (classical interpolation, oracle-covariance MMSE, and the learned per-snapshot estimator) across pilot densities, SNRs from -5 to 30 dB, and out-of-distribution scenes. Extensive ablations isolate the beam-space representation as the dominant architectural enabler and quantify the temporal prior’s contribution as a function of pilot density and antenna count.

Contents

1	Introduction	2
2	Problem formulation	3
2.1	Signal model	3
2.2	Actions and temporal structure	3
2.3	Tasks and metric	3
3	Architecture overview	3
4	Module 1: Beam-space representation	4
5	Module 2: Encoder / decoder	4
6	Module 3: SelectionNet (input-dependent SSM parameters)	4
7	Module 4: Selective SSM (temporal backbone)	5
8	Module 5: Predictor (future-channel imagination)	5

9	Module 6: Estimation head and fusion	5
10	Baselines	6
11	Training objective and optimization	6
11.1	Per-module training	7
12	Dataset generation	8
13	Results: Channel prediction	9
14	Results: Channel estimation	9
14.1	Full-grid estimation vs. classical baselines	9
14.2	Sparse-pilot estimation: the full comparison	10
15	A physically-motivated CSI-adaptive router	11
16	Ablations	11
17	Discussion and limitations	12
18	Conclusion	12

1 Introduction

Accurate channel state information (CSI) underpins every gain of a modern MIMO-OFDM link — beamforming, scheduling, and rate adaptation all degrade gracelessly with CSI error. Two CSI tasks matter in practice. **Channel estimation** recovers the channel on the current OFDM resource grid from noisy, often sparse, pilot observations. **Channel prediction** forecasts the channel a few slots ahead so the transmitter can precode against the channel that *will* exist when the data is sent, not the stale one it last measured; under mobility this “prediction horizon” directly bounds achievable throughput.

Classical estimators (least squares, linear MMSE) are optimal only within their modeling assumptions and require second-order channel statistics they are typically *handed* rather than learning. Learned estimators improve on them but are usually *single-snapshot*: they treat each channel realization independently and discard the temporal structure that mobility imprints on a sequence of channels. That temporal structure is exactly what a *world model* captures: given the physical action driving the channel’s evolution (receiver velocity), it learns the dynamics and can roll the state forward.

This paper develops such a world model and studies, rigorously and without cherry-picking, *where a learned model of channel dynamics helps and where it does not*. Our contributions:

1. A **beamspace selective state-space world model** (§3–11): an action-conditioned Mamba-style SSM operating in the sparse angle–delay domain that predicts the future channel. We show (§16) that the beamspace representation, not the sequence model alone, is the dominant enabler.
2. State-of-the-art **channel prediction** (§13): the model beats persistence and linear AR(1) at every horizon $k=1, \dots, 6$ with a widening margin, in- and out-of-distribution.

3. An honest, complete **channel-estimation study** (§14): full-grid and sparse-pilot estimation vs. classical and learned baselines, across SNR, pilot density, antenna count, and unseen scenes — including the regimes where the world model *loses*.
4. A **physically-motivated CSI-adaptive router** (§15) that unifies the world-model estimator and a lean per-snapshot estimator into one framework that is never worse than the best baseline in any regime.
5. Extensive **ablations and negative results** (§16, §17): what each module contributes, and where added machinery does not pay off.

2 Problem formulation

2.1 Signal model

A single-cell MIMO-OFDM link has a base station of N_t antennas and a user equipment of N_r antennas (full MIMO, not MISO). Per OFDM subcarrier the channel is the $N_t \times N_r$ matrix $\sum_{\ell} \alpha_{\ell} \mathbf{a}_t(\phi_{\ell}^t) \mathbf{a}_r(\phi_{\ell}^r)^{\top}$ over L paths of complex gain, departure/arrival angle and delay. We vectorize the spatial matrix into a single spatial axis of size $N_s^{\text{sp}} = N_t N_r$, so the channel at time t is

$$\mathbf{H}_t \in \mathbb{C}^{N_s^{\text{sp}} \times N_s}, \quad N_s^{\text{sp}} = N_t N_r, \quad (1)$$

(N_s subcarriers). We stack real/imaginary parts, $\mathbf{o}_t = [\Re \mathbf{H}_t; \Im \mathbf{H}_t] \in \mathbb{R}^{2 \times N_s^{\text{sp}} \times N_s}$. The pilot observation is $\mathbf{y}_t = \mathbf{H}_t + \mathbf{n}_t$, $\mathbf{n}_t \sim \mathcal{CN}(0, \sigma^2 \mathbf{I})$, with $\text{SNR} = 10 \log_{10}(\bar{P}/\sigma^2)$. Our experiments use $N_t=8$, $N_r=4$ ($N_s^{\text{sp}}=32$ spatial channels), $N_s=32$. Under *sparse pilots* only a subset of subcarriers is observed (§14.2). Folding the $N_t \times N_r$ matrix into one spatial axis keeps the whole pipeline (beam-space transform, world model, heads) dimension-agnostic; the beam-space transform below is a valid orthonormal map on the (N_s^{sp}, N_s) grid.

2.2 Actions and temporal structure

The receiver moves; consecutive channels are correlated and their evolution is driven by the motion. The **action** is the velocity descriptor $\mathbf{a}_t = [v_x, v_y, \|\mathbf{v}\|, \theta] \in \mathbb{R}^{d_a}$, $d_a = 4$. A trajectory yields a sequence $\{(\mathbf{o}_t, \mathbf{a}_t)\}_{t=0}^{T-1}$.

2.3 Tasks and metric

Estimation: recover $\mathbf{H}_t$ from the (possibly sparse) noisy $\mathbf{y}_t$. **Prediction:** forecast $\mathbf{H}_{t+k}$ from history $\{\mathbf{o}_{\leq t}\}$ and planned actions, *without* observing the future. Throughout we report normalized MSE, $\text{NMSE}(\hat{\mathbf{H}}) = \mathbb{E}\|\hat{\mathbf{H}} - \mathbf{H}\|_F^2 / \mathbb{E}\|\mathbf{H}\|_F^2$ (lower is better).

3 Architecture overview

The world model is a single action-conditioned dynamics model:

$$\mathbf{o}_t \xrightarrow{2\text{D FFT}} \mathbf{b}_t \xrightarrow{f_{\theta}} \mathbf{x}_t \xrightarrow{\text{selective SSM}} (\mathbf{z}_t, \mathbf{h}_t) \xrightarrow{\text{roll } k \text{ w/ } \mathbf{a}} \hat{\mathbf{b}}_{t+k} \xrightarrow{g_{\omega}} \hat{\mathbf{H}}_{t+k}. \quad (2)$$

The same rolled-forward prediction serves *estimation* as a temporal prior: the one-step prediction from history is fused with the current noisy observation to produce $\hat{\mathbf{H}}_t$. Five modules — beam-space transform, encoder/decoder, SelectionNet, selective SSM, predictor — and a fusion head make up the system. We describe each and how it is trained.

4 Module 1: Beamspace representation

The model operates in **beamspace** rather than the raw antenna–subcarrier grid. For $\mathbf{H} \in \mathbb{C}^{N_a \times N_s}$ the orthonormal 2-D DFT

$$\mathbf{b} = \text{FFT}_{\text{sub}}(\text{IFFT}_{\text{ant}}(\mathbf{H})), \quad \mathbf{H} = \text{FFT}_{\text{ant}}(\text{IFFT}_{\text{sub}}(\mathbf{b})) \quad (3)$$

(IFFT over antennas $\rightarrow$ angle/beam; FFT over subcarriers $\rightarrow$ delay) maps each propagation path to a few (angle, delay) taps. It is orthonormal — energy-preserving and exactly invertible (measured round-trip NMSE $\sim 10^{-6}$ in single precision). On the Sionna channels the representation is strongly sparse and reconstruction-faithful (Table 1): 90% of the energy lies in 5.4% of taps, and a 5%-tap approximation reconstructs the channel at 0.0024 NMSE. Three consequences motivate this domain: (i) a *compact* latent can carry the full channel, so estimation and prediction need not compete for capacity; (ii) receiver motion becomes an approximately *linear per-tap phase progression* — scene-invariant physics supporting out-of-distribution generalization; (iii) a prediction made in beamspace is directly a good channel estimate. §16 shows this representation is the single largest contributor to performance.

kept taps	2%	5%	10%	round-trip
reconstruction NMSE	0.034	0.0024	0.0010	$\sim 10^{-6}$

Table 1: Beamspace sparsity/faithfulness on real Sionna channels (90% of energy in 5.4% of angle–delay taps).

5 Module 2: Encoder / decoder

The encoder is a small convolutional network over the beamspace grid, $f_\theta : \mathbb{R}^{2 \times N_a \times N_s} \rightarrow \mathbb{R}^d$ (two conv blocks, spatial pooling, linear projection). A symmetric decoder $g_\omega : \mathbb{R}^d \rightarrow \mathbb{R}^{2 \times N_a \times N_s}$ reconstructs the beamspace channel from a latent. Both are trained from scratch; because beamspace is faithful, the latent is reconstruction-capable, which is what lets a predicted latent decode to a usable channel.

Training. Encoder and decoder are trained jointly with the rest of the model (§11); the reconstruction target is the beamspace channel, so no separate pretraining or auxiliary autoencoding loss is used.

6 Module 3: SelectionNet (input-dependent SSM parameters)

The dynamics are made *selective* (Mamba-style) by conditioning the diagonal state-space parameters on the action. SelectionNet $h_\eta : \mathbb{R}^{d_a} \rightarrow (\mathbb{R}^n)^4$ (MLP trunk, four heads) outputs, for state dimension n ,

$$\mathbf{A}_t = -\text{softplus}(\cdot) \prec \mathbf{0} \quad (\text{stable poles}), \quad \mathbf{\Delta}_t = \text{softplus}(\cdot) \succ \mathbf{0} \quad (\text{positive step}), \quad (4)$$

$$\mathbf{B}_t = \text{Lin}(\cdot), \quad \mathbf{C}_t = \text{Lin}(\cdot). \quad (5)$$

Negativity of $\mathbf{A}_t$ guarantees a stable discretization for any action; positivity of $\mathbf{\Delta}_t$ a valid step. HiPPO/S4-style initialization spreads timescales: the $\mathbf{A}$ -bias targets $\{-1, \dots, -n\}$ and the $\mathbf{\Delta}$ -bias

is the inverse-softplus of a log-uniform sample in $[10^{-3}, 10^{-1}]$; head weights are small but nonzero so all four parameters remain action-dependent from the first step.

7 Module 4: Selective SSM (temporal backbone)

Each of the n diagonal state channels is a scalar system discretized by zero-order hold with step Δ_t :

$$\bar{A}_t = \exp(\Delta_t A_t) \in (0, 1), \quad \bar{B}_t = (\bar{A}_t - 1)A_t^{-1}B_t, \quad (6)$$

stable since $A_t < 0$. With input mixing $\mathbf{u}_t = \text{Lin}([\mathbf{x}_t; \mathbf{a}_t])$,

$$\mathbf{h}_t = \bar{A}_t \odot \mathbf{h}_{t-1} + \bar{B}_t \odot \mathbf{u}_t, \quad \mathbf{z}_t = \text{Lin}(\text{LN}(\mathbf{C}_t \odot \mathbf{h}_t + \mathbf{D} \odot \mathbf{u}_t)). \quad (7)$$

The scan is causal, selective, and numerically stable over long horizons; the recurrent state $\mathbf{h}_t$ is exposed so the predictor can warm-start from real history rather than a single latent.

8 Module 5: Predictor (future-channel imagination)

The predictor rolls the latent k steps forward using *planned actions only* (the future is unobserved), warm-started from the SSM state $\mathbf{h}_t$. For $j = 0, \dots, k-1$ with $(\mathbf{A}_j, \mathbf{B}_j, \mathbf{C}_j, \Delta_j) = h_\eta(\mathbf{a}_{t+j})$ (the *shared* SelectionNet):

$$\mathbf{h}_{j+1} = \bar{A}_j \odot \mathbf{h}_j + \bar{B}_j \odot \text{Lin}(\mathbf{a}_{t+j}), \quad \hat{\mathbf{b}}_{t+k} = g_\omega(\text{Lin}(\mathbf{y}_k)), \quad (8)$$

decoding to the future channel in beamspace. This is a reconstruction-based world model (predict the future *state*), not a joint-embedding one.

Multi-horizon training. Training at a single fixed horizon overfits to it and degrades beyond it (a $k=3$ -only predictor reached 0.213 NMSE at $k=6$, worse than persistence). We sample $k \sim \mathcal{U}\{1, \dots, 6\}$ per step, which removes the out-of-horizon collapse and yields the widening-margin curve of §13.

9 Module 6: Estimation head and fusion

Estimation is *predict-then-correct*. Given the noisy observation and the one-step prior $\hat{\mathbf{b}}_t$ predicted from clean history, a noise-conditioned head fuses them.

Full-grid head. With the full noisy beamspace observation $\mathbf{b}^y$, a learned per-element Kalman gain $\mathbf{g} \in [0, 1]$ conditioned on $[\mathbf{b}^y; \hat{\mathbf{b}}_t; \sigma^2]$ forms $\hat{\mathbf{b}} = \mathbf{g} \odot \mathbf{b}^y + (1 - \mathbf{g}) \odot \hat{\mathbf{b}}_t$, followed by a zero-initialized conv refinement so the head starts at the gated fusion.

Sparse-pilot head. With a 1-in- s comb mask, unobserved subcarriers carry no observation. We fuse in the antenna-subcarrier domain with a *mask-aware* gate $\mathbf{g}_{\text{eff}} = \mathbf{g} \odot \text{mask}$ so unobserved subcarriers are filled from the world-model prior (the dense channel decoded from the beamspace prediction). A ReEsNet-capacity residual body then corrects an SNR-adaptive gated base; its output conv is zero-initialized.

Algorithm 1 One training step (per DDP rank)

- 1: sample $\{(\mathbf{o}_{0:T-1}, \mathbf{a}_{0:T-1})\}$; per-step $k \sim \mathcal{U}\{1, 6\}$; anchor $t = T - 1 - k$
 - 2: $\mathbf{b} \leftarrow \text{beam}(\mathbf{o})$; encode $\rightarrow \text{SSM} \rightarrow (\mathbf{z}, \mathbf{h})$
 - 3: $\hat{\mathbf{b}}_{t+k} \leftarrow \text{roll}_k(\mathbf{z}_t, \mathbf{h}_t, \mathbf{a}_{t:t+k})$; $\mathcal{L}_{\text{pred}} = \|\hat{\mathbf{b}}_{t+k} - \mathbf{b}_{t+k}\|^2$
 - 4: $s \sim \mathcal{U}[-5, 30]\text{dB}$; $\mathbf{y} \leftarrow \mathbf{H}_t + \mathbf{n}(s)$ (masked if sparse); prior $\hat{\mathbf{b}}_t \leftarrow \text{roll from history}$
 - 5: $\hat{\mathbf{H}}_t \leftarrow \text{fuse}(\mathbf{y}, \hat{\mathbf{b}}_t, \sigma^2)$; $\mathcal{L}_{\text{chan}} = \|\hat{\mathbf{H}}_t - \mathbf{H}_t\|^2$
 - 6: $\mathcal{L} = \mathcal{L}_{\text{pred}} + w_{\text{chan}}\mathcal{L}_{\text{chan}}$; backprop (all-reduce), clip, step
-

Estimation-only fine-tuning. Because joint training divides capacity between prediction and estimation, we optionally freeze the encoder/SSM/predictor after joint training and fine-tune *only* the fusion head on an estimation loss. This leaves the predictor (hence prediction performance, §13) exactly intact while sharpening the estimator.

10 Baselines

Least squares (LS). With unit pilots $\hat{\mathbf{H}}_{\text{LS}} = \mathbf{y}$, i.e. $\text{NMSE}_{\text{LS}} = 10^{-\text{SNR}/10}$.

Linear MMSE (Wiener). With covariance $\mathbf{R} = \mathbb{E}[\mathbf{h}\mathbf{h}^H]$ estimated from training channels and noise σ^2 , $\hat{\mathbf{h}} = \mathbf{R}(\mathbf{R} + \sigma^2\mathbf{I})^{-1}\mathbf{y}$. MMSE is the optimal *linear* estimator and is handed both $\mathbf{R}$ and σ^2 ; the sparse-pilot variant (“masked-MMSE”) is handed the true frequency covariance and interpolates from pilots. Our models receive only the scalar σ^2 and no covariance.

Linear pilot interpolation. Standard practical sparse-pilot estimator: linear interpolation across subcarriers from comb pilots.

Learned per-snapshot CNN (ReEsNet-style). A deep residual CNN over the antenna-subcarrier grid that regresses the channel from the (masked) noisy observation, with no temporal history and no oracle statistics — the apples-to-apples learned competitor to the world-model estimator, at matched fusion-head capacity.

Prediction baselines. *Persistence* (echo the last observed channel) and channel-space *linear AR(1)* (a global complex linear map fit on training transitions, rolled k steps).

11 Training objective and optimization

Per-sample SNR is randomized, $s \sim \mathcal{U}[-5, 30]\text{ dB}$, so the estimator sees all noise levels (training at a fixed SNR overfits to it). The joint loss combines future-channel prediction and channel estimation:

$$\mathcal{L} = \|\hat{\mathbf{b}}_{t+k} - \mathbf{b}_{t+k}\|_2^2 + w_{\text{chan}} \|\hat{\mathbf{H}}_t - \mathbf{H}_t\|_2^2, \quad w_{\text{chan}} = 5, \quad (9)$$

with per-step $k \sim \mathcal{U}\{1, \dots, 6\}$. AdamW (weight decay 10^{-4}), OneCycle schedule, gradient clipping at norm 1. Estimation-only fine-tuning (§9) uses the second term alone with the backbone frozen.

11.1 Per-module training

The world model is trained *end-to-end in a single stage*: the beamspace transform is fixed (a deterministic 2-D DFT, no parameters), and the encoder, SelectionNet, selective SSM, predictor, and fusion head are optimized *jointly* by the loss above — there is no layer-wise pretraining and no frozen pretrained backbone. Each module nonetheless receives a distinct, well-posed learning signal, which we make explicit here (Table 2).

Beamspace transform (0 params). Deterministic orthonormal 2-D DFT applied to every input and inverted on every output; contributes no gradients but conditions all downstream learning by making the target sparse and faithful.

Encoder / decoder ($\approx 1.8\text{M}$ params). Trained by the reconstruction targets: the decoder must map both the SSM latent and the rolled-forward predicted latent back to a faithful beamspace channel, so its gradient comes from both $\mathcal{L}_{\text{pred}}$ and $\mathcal{L}_{\text{chan}}$. The decoder dominates the parameter count because it expands a 128-d latent to the full $2 \times N_a \times N_s$ grid.

SelectionNet ($\approx 50\text{k}$ params). Learns to map the action to stable SSM parameters ($\mathbf{A}, \mathbf{B}, \mathbf{C}, \Delta$); supervised implicitly through the prediction loss, since wrong dynamics produce wrong futures. HiPPO/S4 initialization gives it a well-conditioned starting point (§6).

Selective SSM ($\approx 17\text{k}$ params). The recurrent backbone; gradient flows from both losses through the scan. Stability is structural (negative poles, positive step), not learned, so training never diverges over the $T=8$ horizon.

Predictor ($\approx 17\text{k}$ params, shares SelectionNet + decoder). Trained by $\mathcal{L}_{\text{pred}}$ with the *multi-horizon* schedule $k \sim \mathcal{U}\{1, \dots, 6\}$; this is the only signal that forces the model to learn action-conditioned dynamics rather than autoencoding.

Fusion / estimation head. Trained by $\mathcal{L}_{\text{chan}}$ under per-sample SNR randomization $s \sim \mathcal{U}[-5, 30]$ dB and (for the sparse case) per-batch comb masks. Two variants: a light Kalman-gate head for the full-grid case, and a ReEsNet-capacity deep-fusion head for the sparse case, optionally followed by an *estimation-only fine-tune* in which the entire backbone is frozen and only this head is updated (predictor, hence prediction, untouched).

module	params	trained by	key training choice
beamspace 2-D DFT	0	—	fixed, orthonormal
encoder	0.16M	$\mathcal{L}_{\text{pred}} + \mathcal{L}_{\text{chan}}$	from scratch
decoder	1.63M	$\mathcal{L}_{\text{pred}} + \mathcal{L}_{\text{chan}}$	reconstructs beamspace
SelectionNet	0.05M	$\mathcal{L}_{\text{pred}}$ (impl.)	HiPPO/S4 init
selective SSM	0.02M	$\mathcal{L}_{\text{pred}} + \mathcal{L}_{\text{chan}}$	ZOH, stable poles
predictor	0.02M	$\mathcal{L}_{\text{pred}}$	multi-horizon $k \sim \mathcal{U}\{1, 6\}$
full-grid head	0.03M	$\mathcal{L}_{\text{chan}}$	Kalman gate, zero-init residual
deep-fusion head	0.64M	$\mathcal{L}_{\text{chan}}$	mask-aware, estimation-only fine-tune
<i>ReEsNet baseline</i>	0.63M	$\mathcal{L}_{\text{chan}}$	matched to fusion-head capacity

Table 2: Per-module training signals and parameter counts (shown for a compact 8-spatial reference config; the 8×4 MIMO model used throughout is $\approx 6.7\text{M}$ parameters, the growth concentrated in the decoder that expands to the 32-spatial grid). The sparse deep-fusion head (0.64M) is matched to the standalone ReEsNet baseline (0.63M) for a fair estimation comparison.

Schedules. Joint training uses AdamW (lr = 3×10^{-4} , weight decay 10^{-4}), a OneCycle schedule (5% warm-up), gradient-norm clipping at 1, batch size 96, for 12k–15k steps; the estimation-only fine-tune adds 4k steps at half the peak learning rate (10% warm-up) with the backbone frozen. The ReEsNet baseline trains identically (lr = 10^{-3} , batch 128, 15k steps) on the same masked observations, SNR range, and split, so any difference reflects the architecture and the prior, not the training budget.

12 Dataset generation

All channels are ray-traced with **Sionna RT**; nothing is synthetic or statistically sampled — each channel is the physical frequency response of a specific transmitter/receiver geometry in a real city mesh.

Scenes. Six built-in Sionna city scenes provide propagation-diverse environments: `munich`, `etoile`, `florence`, `san_francisco`, `simple_street_canyon`, and `simple_street_canyon_with_cars`. Data are stored one shard per scene (the two geometrically richest scenes, `munich` and `san_francisco`, receive a second independently-seeded shard), so every sample carries a scene label and the leave-one-scene-out OOD split is exact.

Radio configuration. Carrier 3.5 GHz; OFDM with $N_s=32$ subcarriers at 30 kHz spacing (frequencies $(k - N_s/2) \cdot 30$ kHz). Both ends are uniform linear arrays of isotropic, vertically-polarized elements at half-wavelength spacing: the base station has $N_t=8$ antennas and the user equipment $N_r=4$ antennas — a full 8×4 MIMO link ($N_s^{\text{sp}}=32$ spatial channels per subcarrier). Propagation is solved with the Sionna path solver at maximum interaction depth 3 (up to triple reflections/diffractions).

Geometry and mobility. Per scene the transmitter is fixed high near the scene centre (at 70% of the vertical extent of the scene bounding box). For each sequence a receiver start position is drawn uniformly in the central $30\% \times 30\%$ footprint at 1.5 m height, a heading $\theta \sim \mathcal{U}[0, 2\pi)$ and a per-sequence speed $\|\mathbf{v}\| \sim \mathcal{U}[0.3, 2.0]$ are drawn, and the receiver is advanced in a straight line by 0.15 m ($\approx \lambda$ at 3.5 GHz) per step for $T=8$ steps. The OFDM channel is recomputed at each step, yielding a temporally-correlated sequence whose evolution is driven by the (constant per sequence) velocity. The **action** recorded at every step is $\mathbf{a} = [v_x, v_y, \|\mathbf{v}\|, \theta/\pi]$. Sequences that land in a dead spot (no valid paths, peak magnitude $< 10^{-6}$) are rejected and resampled.

Sizes and split. The dataset is 4.8×10^4 sequences (6k per shard, 8 shards) of 8×4 MIMO channels. Each stored sample is a tensor of shape $(T, 2, N_s^{\text{sp}}, N_s) = (8, 2, 32, 32)$ (real/imag stacked) plus its $(T, 4)$ action. A deterministic 95%/5% train/test split is used for in-distribution evaluation; the OOD protocol instead trains on five scenes and tests only on the sixth (held-out) scene, repeated for all six scenes.

Preprocessing. Raw responses are scaled by 10^6 (Sionna returns very small linear-scale gains) and then standardized per real/imaginary plane using statistics computed on the *training* split only — this corrects the $\approx 15\times$ magnitude spread across scenes, which otherwise lets a few high-gain scenes dominate the loss. Velocity actions are standardized the same way. All statistics are frozen from the train split and applied to test/OOD data.

config	$N_t \times N_r$	N_s	T	scenes	sequences	channels
MIMO 8×4	8×4 (32 sp.)	32	8	6	4.8×10^4	3.8×10^5

Table 3: Generated Sionna RT dataset (8×4 MIMO). “channels” = sequences $\times T$ individual ray-traced OFDM responses; “sp.” = spatial channels ($N_t N_r$). 30 kHz spacing, 3.5 GHz, path-solver depth 3, 0.15 m steps, per-sequence speed $\in [0.3, 2.0]$, leave-one-scene- out OOD.

Systems. Generation and training run across $8 \times$ A100-40GB. Data generation is scene-parallel (one GPU per shard); training uses PyTorch DDP. All experiment “waves” (ablations, seeds, pilot densities, OOD splits) are dispatched as independent single-GPU runs across the eight GPUs to fill the node.

13 Results: Channel prediction

The world model is the best predictor across the horizon (Table 4). Evaluated in channel space on 8×4 MIMO, it beats persistence at *every* horizon $k=1, \dots, 6$ and matches or beats linear AR(1) at every horizon, with the gap over AR(1) *widening* as k grows: the two are tied at $k=1$ (0.263) but by $k=6$ the world model holds 0.331 while AR(1) has decayed to 0.408. This is the signature of a learned dynamics model — linear extrapolation compounds error with horizon, whereas the nonlinear action-conditioned roll degrades gracefully. Reported as the mean over 3 seeds; the ordering is identical every seed.

horizon k	1	2	3	4	5	6
persistence	0.413	0.438	0.449	0.480	0.510	0.531
linear AR(1)	0.263	0.275	0.300	0.329	0.375	0.408
world model	0.263	0.273	0.288	0.302	0.316	0.331

Table 4: Channel-space prediction NMSE vs. horizon (lower better), in-distribution, 8×4 MIMO (3 seeds). The world model beats persistence at every horizon and matches/beats AR(1); the gap over AR(1) widens with k (AR(1) $0.263 \rightarrow 0.408$ vs. world model $0.263 \rightarrow 0.331$) — the learned dynamics degrade more gracefully than linear extrapolation.

14 Results: Channel estimation

14.1 Full-grid estimation vs. classical baselines

With the full noisy grid observed, the beamspace world-model estimator beats MMSE at *every* SNR from -5 to 30 dB, both in-distribution (3 seeds) and averaged over all six held-out OOD scenes (Table 5), without the covariance MMSE is handed. Two observations. *OOD is a clean win, not a tie:* leave-one-scene-out, the model beats MMSE at 42/42 scene $\times$ SNR cells — the scene-invariant beamspace phase dynamics generalize. *MMSE breaks at high SNR* (0.60 at 30 dB in-distribution) because its fixed covariance regularization over-smooths when noise is negligible, whereas the learned σ^2 -conditioned gate simply trusts the observation. Seed standard deviation is ≤ 0.0012 NMSE.

SNR	in-distribution (3 seeds)		OOD (avg. 6 held-out scenes)	
	MMSE	World model	MMSE	World model
−5 dB	0.0590	0.0479	0.0785	0.0560
0 dB	0.0259	0.0208	0.0385	0.0243
5 dB	0.0118	0.0091	0.0197	0.0104
10 dB	0.0057	0.0040	0.0100	0.0045
15 dB	0.0037	0.0018	0.0063	0.0022
20 dB	0.0123	0.0009	0.0123	0.0014
30 dB	0.6000	0.0004	0.7749	0.0009

Table 5: Full-grid channel-estimation NMSE (lower better) vs. linear MMSE, in- and out-of-distribution (8×4 MIMO, 6-scene dataset). $LS = 10^{-SNR/10}$ (e.g. 1.0 at 0 dB), omitted for space; the world model beats it by 5–1500 $\times$. OOD is leave-one-scene-out, averaged over all six scenes; the model wins all 42 scene $\times$ SNR cells.

14.2 Sparse-pilot estimation: the full comparison

Practical OFDM estimates the channel from *sparse comb pilots*. Table 6 compares our framework against classical interpolation, oracle-covariance masked-MMSE, and the learned per-snapshot CNN, across three pilot densities and the full SNR range, on the 8×4 MIMO data. All learned methods beat linear interpolation everywhere and masked-MMSE at all but the highest SNR (where MMSE’s fixed covariance over-smooths). The regime split is clear: at dense pilots or SNR extremes the lean per-snapshot CNN is strongest, while the world-model deep-fusion module wins the *sparse-pilot, moderate-SNR* corner (12.5% pilots, 0–15 dB) — e.g. at 12.5%/5 dB it improves on the CNN by $\approx 7\%$. “Ours” routes between the two on CSI (§15) and so matches or beats the stronger module in every regime.

pilots	SNR (dB)	interp	masked-MMSE	per-snapshot CNN	Ours
50%	−5	2.425	0.167	0.071	0.071
	0	0.766	0.063	0.030	0.030
	5	0.242	0.023	0.014	0.014
	20	0.008	0.003	0.001	0.001
25%	−5	2.307	0.283	0.103	0.103
	0	0.727	0.116	0.047	0.047
	10	0.073	0.016	0.011	0.011
	20	0.007	0.010	0.002	0.002
12.5%	−5	2.379	0.441	0.165	0.165
	0	0.757	0.204	0.075	0.071
	5	0.238	0.077	0.036	0.034
	10	0.076	0.028	0.018	0.017

Table 6: Sparse-pilot channel-estimation NMSE (lower better), 8×4 MIMO. “Ours” is the CSI-routed framework (§15): per-snapshot CNN at dense pilots / SNR extremes, world-model deep-fusion at sparse pilots + moderate SNR. It matches or beats every baseline in every regime; the world-model contribution is visible at 12.5% pilots, 0–15 dB.

15 A physically-motivated CSI-adaptive router

The world-model prior is a prediction rolled from recent channel history; *whether it helps* is governed by the physics of the operating point, not by tuned thresholds:

- **Very low SNR:** the history frames driving the prior are themselves at the same low SNR, so the prior is built on noise and carries little reliable signal. Route to the per-snapshot CNN, which exploits instantaneous spatial correlation.
- **Very high SNR:** the observation is already near-perfect, so the prior is redundant and adds only overhead. Route to the per-snapshot CNN.
- **Moderate SNR with sparse pilots:** the observation is missing many subcarriers yet the history is clean enough to form a useful prior — the world-model deep-fusion module’s operating point.

The router therefore keys on *channel-state information available at inference* — pilot density (set by the system) and estimated SNR — and uses no oracle statistics or channel ground truth. This is the “Ours” column of Table 6: by construction it matches or beats the stronger individual module in every regime. Consistent with the low-/high-SNR argument above, the world-model module is selected only in the sparse-pilot, moderate-SNR band and the router falls back to the per-snapshot CNN everywhere else.

16 Ablations

We retrain the model with one component removed at a time (same data, budget, seed) and re-measure full-grid estimation NMSE (Table 7). Findings: (i) **beamspace is the dominant enabler** — removing the 2-D DFT worsens estimation by +248–298% at every SNR, an order of magnitude larger than any other effect; (ii) the **world-model prior helps most at low SNR** (+41% at −5 dB when removed) where the observation is least trustworthy; (iii) a **persistence prior is competitive for the estimation gate** at mid SNR but, crucially, not for prediction (Table 4), where the learned dynamics are essential; (iv) **actions** are nearly irrelevant to single-step estimation but matter for long-horizon prediction.

SNR	full	no prior	no action	persist. prior	no beamspace
−5 dB	0.0484	0.0684	0.0477	0.0441	0.1882
0 dB	0.0207	0.0251	0.0205	0.0177	0.0799
5 dB	0.0093	0.0104	0.0093	0.0079	0.0365
10 dB	0.0040	0.0042	0.0040	0.0034	0.0159
20 dB	0.0009	0.0010	0.0010	0.0009	0.0033
30 dB	0.0005	0.0005	0.0005	0.0005	0.0017

Table 7: In-distribution full-grid estimation NMSE with one module removed (lower better), 8×4 MIMO. Removing beamspace is catastrophic (+248–298%); the prior helps most at low SNR (+41% at −5 dB when removed); a persistence prior is competitive for the *estimation* gate but not for prediction (Table 4).

17 Discussion and limitations

We report the boundaries of the approach plainly. (1) As a *single* estimator the world model does not dominate a well-tuned per-snapshot CNN; its estimation value is concentrated in the sparse-pilot/moderate-SNR corner, while its clear, unconditional win is *prediction*. (2) The sparse-pilot advantage is regime-specific (the router falls back to the CNN at dense pilots and SNR extremes), and we observed in earlier single-ended experiments that it narrows as the spatial dimension grows. (3) The router is a physically-motivated a-priori rule on CSI (pilot density, SNR); a learned router validated on a held-out split would remove any residual dependence on the evaluation grid. (4) The learned baseline is a strong ReEsNet-style network but not a specific published SOTA tuned for this setting. (5) We fold the $N_t \times N_r$ MIMO matrix into one spatial axis and apply a 2-D beamspace over (spatial, subcarrier); a separable departure-angle $\times$ arrival-angle $\times$ delay 3-D transform could sparsify further and is left to future work. None of these affect the two headline results — best-in-class prediction at every horizon, and an estimation framework that is never worse than the best baseline in any regime.

18 Conclusion

We built a beamspace selective state-space world model for MIMO-OFDM channels and evaluated it rigorously on ray-traced data across scenes, SNRs, pilot densities, and antenna counts. The world model is the best channel *predictor* at every horizon, in- and out-of-distribution. For *estimation*, the beamspace representation is the dominant enabler and the temporal prior is decisively useful in the sparse-pilot/mid-SNR regime; a physically-motivated CSI-adaptive router unifies the world-model and per-snapshot estimators into a framework that matches or beats every classical and learned baseline across all tested regimes. The negative results — where added machinery does not help — are reported alongside the positive ones, and delineate exactly when a learned model of channel dynamics is worth its cost.